About the Organization

Redhorse transforms the way government uses data and technology. To support this mission, we are seeking a **Machine Learning Engineer** to support the Office of the CTO. In this role, you will help develop cutting-edge capabilities that can be utilized across our diverse customer set, enhancing our ability to leverage generative AI technologies for mission-focused government applications.

Responsibilities include:

- Design, develop, and maintain machine learning models with a focus on generative AI, including but not limited to natural language processing and computer vision.
- Engage in prompt engineering, exploring innovative ways to interact with large language models (LLMs) to achieve desired outputs.
- Integrating open-source machine learning projects and the latest advancements into our solutions.
- Develop and implement strategies for fine-tuning large language models for specific tasks or datasets, enhancing model performance and accuracy.
- Work closely with the product and engineering teams to deploy scalable models.
- Lead cross-functional initiatives to pilot and scale AI technologies across the organization, aligning with Redhorse's mission and strategic objectives.
- Develop prototypes to demonstrate to clients the value and impact of Generative AI capabilities.

Basic Requirements

- 3 years of experience in machine learning, data science, or a related role, with a focus on generative AI and natural language processing.
- Proficiency in Python, and familiarity with machine learning frameworks like PyTorch and LangChain.
- Experience with prompt engineering, model fine-tuning, and deploying machine learning models .
- Demonstrated ability to work on complex problems and translate requirements into robust AI solutions.

Desired Experience:

We encourage all candidates who meet the basic requirements to apply, even if you do not have any of the desired experience. Join us at Redhorse to play a pivotal role in harnessing the power of AI to transform government technology and data practices.

- Degree in Computer Science, Engineering, Mathematics, or a related field.
- Contributions to open-source machine learning projects or active participation in the AI research community.
- Experience with cloud computing platforms (AWS, Google Cloud, Azure) and containerization technologies (Docker, Kubernetes).
- Familiarity with the latest trends in AI ethics and responsible AI practices.
- Demonstrating prototypes to clients.